

## DAT 103

### Datamaskiner og operativsystemer (Computers and Operating Systems)

#### Supplementary exercises (Set 3)

##### Problem 1

Consider an operating system running on a computer system with 16 cores uses the *many-to-many* model to map user threads to kernel threads. If  $N$  user threads are created for a particular process and  $N > 16$ , how many kernel threads will be allocated to this process?

**Solution.**

At most  $N$

##### Problem 2

Describe the differences among short-term, and long-term scheduling.

**Solution.**

Short-term (CPU scheduler) selects from jobs in memory those jobs that are ready to execute and allocates the CPU to them.

Long-term (job scheduler) determines which jobs are brought into memory for processing.

The primary difference is in the frequency of their execution. The short-term must select a new process quite often. Long-term is used much less often since it handles placing jobs in the system and may wait a while for a job to finish before it admits another one

##### Problem 3

Describe the actions taken by a kernel to context-switch between processes.

**Solution.**

In general, the operating system must save the state of the currently running process and restore the state of the process scheduled to be run next. Saving the state of a process typically includes the values of all the CPU registers in addition to memory allocation. Context switches must also perform many architecture-specific operations, including flushing data and instruction caches.

##### Problem 4

What resources are used when a thread is created? How do they differ from those used when a process is created?

**Solution.**

Because a thread is smaller than a process, thread creation typically uses fewer resources than process creation. The latter requires allocating a process control block (PCB), a rather large data structure, while the former involves allocating a small data structure to hold a register set, stack, and priority.

##### Problem 5

What advantage is there in having different time-quantum sizes at different levels of a multilevel queueing system?

**Solution.**

Processes that need more frequent servicing – for instance, interactive processes such as editors – can be in a

queue with a small time quantum; otherwise, processes can be in a queue with a larger quantum, requiring fewer context switches to complete the processing and thus making more efficient use of the computer.

### Problem 6

Explain the how the following scheduling algorithms discriminate either in favor of or against short processes:

- (a) FCFS
- (b) RR
- (c) Multilevel feedback queues

**Solution.**

- (a) against short jobs, since any short job arriving after a long job will have a longer waiting time.
- (b) treats all jobs equally, so short jobs will be able to leave the system faster, since they will finish first.
- (c) work similarly to the RR algorithm

### Problem 7


Consider the following table of arrival time and burst time for three processes  $P_1$ ,  $P_2$  and  $P_3$ :

Process	Arrival Time	Burst Time
$P_1$	0	9
$P_2$	1	4
$P_3$	2	9

The shortest-remaining-time-first (i.e., preemptive shortest job first) scheduling algorithm is used. Scheduling is carried out only at arrival or completion of processes.

- (a) Draw the Gantt Chart of the execution.

**Solution.**



- (b) What is the average waiting time for the three processes?

**Solution.**

5

Process  $P_1$  is allocated processor at 0 ms as there is no other process in ready queue.  $P_1$  is preempted after 1 ms as  $P_2$  arrives at 1 ms and burst time for  $P_2$  is less than remaining time of  $P_1$ .  $P_2$  runs for 4ms.  $P_3$  arrived at 2 ms but  $P_2$  continued as burst time of  $P_3$  is longer than  $P_2$ . After  $P_2$  completes,  $P_1$  is scheduled again as the remaining time for  $P_1$  is less than the burst time of  $P_3$ .  $P_1$  waits for 4 ms,  $P_2$  waits for 0 ms and  $P_3$  waits for 11 ms. So average waiting time is  $(0+4+11)/3 = 5$ .

## Problem 8

Consider the following table of processing time and period for two periodic processes  $P_1$  and  $P_2$ :


Process	Processing time	Period
$P_1$	5	8
$P_2$	3	5

The completion deadlines of each process are the beginning of its next period. For example, the completion deadlines of  $P_1$  are at time 8, 16, 24, ..., while the completion deadlines of  $P_2$  are at time 5, 10, 15, 20, ...

- 8.1 Use *Rate-Monotonic Scheduling* to schedule the two processes, where the priority is assigned based on the inverse of each process' period. At  $time = 20$ , can the two processes meet all their deadlines? Draw the Gantt Chart of the execution to justify your answer.

**Solution.**


*Rate-Monotonic Scheduling* schedules periodic tasks using a **static** priority policy with preemption, i.e., the priority does not change over time. Since the priority is assigned based on the inverse of each process' period, that means  **$P_2$  has higher priority than  $P_1$** , and this priority does not change. In the following gantt charts,  $P_{i-j}$  is used to indicate the  $j$ -th occurrence of  $P_i$ . As you can see,  $P_2$  always **meets** its deadlines, while  $P_1$  always **misses** its deadlines.



- 8.2 Use *Earliest Deadline First* to schedule the two processes. At  $time = 20$ , can the two processes meet all their deadlines? Draw the Gantt Chart of the execution to justify your answer.

**Solution.**

*Earliest Deadline First* schedules periodic tasks using a **dynamic** priority policy with preemption, i.e., the priority *does change* over time. In this subquestion,  $P_2$  has higher priority than  $P_1$  in the beginning, however, it will change dynamically over time. From the gantt chart below, both  $P_1$  and  $P_2$  meet their first deadline, but have problem in keeping the rest. Note that  $P_{2-3}$  is scheduled prior to  $P_{1-2}$  is due to the former has an earlier deadline.



**Hint:** Both Rate-Monotonic Scheduling and Earliest Deadline First are priority scheduling with preemption.